

Web Development and API Design

Lesson 12:

Online Multi-Player Game

Prof. Andrea Arcuri

Goals


- See a full working example of a non-trivial, multi-player, online game based on what learned so far in class
- Learn how to deploy your app on a cloud provider
- Some discussions on databases...

Online Game


Multi-player Connect4

- Turn-based, but need *WebSockets* to get informed when opponent has done his/her move
- How to pair players on different machines?
- How to prevent cheating?
- etc.

- REST API call to start match, same endpoint (IP:port)
- Put on queue until at least 2 players



- When match starts, **each player** creates a *WebSocket* toward the server
- To avoid cheating, the **state** of the match is on server, the players only see a **view** of it (recall users can do whatever on his/her browser...)
- Match state: **board** plus 2 *WebSockets* for the 2 players
- At each state change on server, opponent gets informed via *WebSocket*
- Actions out of sequence must be discarded (possible cheater)



WebSocket Security

- Assume match with id **X** between players **A** and **B**
- How does server know that incoming *WebSocket* request for player **A** on match **X** is actually made by player **A**???
- At least 2 options
 1. As first WS message is in HTTP, can use same auth
 2. Create a **unique, one-time-use random token** in a **secured** way for **A** (eg authed REST endpoint), associate such token to **A** on server, auth the WS when such token is provided
 - This is a more general approach, not tied to implementation details of how the WS is established
 - Recall that WS is for a single user (own TCP socket), so need auth only once, and not like on every single request like in HTTP (eg using cookies)

WebSocket Cost

- WSs are **EXPENSIVE**, as TCP sockets are OS resources
- Need to keep open a TCP socket **for each** WS
- In HTTP, if running out of ports, can close TCP after resolving each incoming request
 - Recall a TCP connection is defined by 4 coordinates: IP and port of client, and IP and port of server
 - Ie, in HTTP keeping a TCP on after a request is only for performance reasons, eg if expecting other following requests
 - See also: “*Connection: Keep-Alive*” HTTP header

WebSocket Topics

- There can be different types of communications between a client and server over WS
 - eg many different functionalities and operations
- Too expensive to open a different TCP socket for each operation
- **Must re-use same WS** for each different kind of operation
- (Simple) Solution: wrap JSON data into an object defining a discriminating “*topic*” field, which identifies the operation
 - Note: there are more sophisticated message protocols like STOMP
- Server will decide what to execute based on topic’s field value

Example

```
{  
  "topic": "..."  
  "data": {...}  
}
```

```
{  
  "topic": "update"  
  "data": {  
    "matchId": 42,  
    "boardDto": {...},  
    "isX": true,  
    "opponentId": 1234  
  }  
}
```

Cloud Deployment

Cloud Deployment

- Different companies provide cloud hosting solutions for your applications, which frees you from hardware issues, but for a price
- *Amazon Web Services (AWS)* is perhaps the most famous/used one
 - eg, *Netflix* runs on AWS
- *Automated scaling*: if you need more load, automatically rent more nodes, and automatically scale down if less load
 - this is also good for applications targeting a specific country (eg Norway), in which you will not get much load during the night

Definition of “Cloud”



Heroku

- One of the main cloud providers
- At the time of this writing, it provides *easy* to use *free* hosting
 - note, this might change at any time
- Supporting NodeJS applications
 - and many others

Using Heroku

- First you need to create an account at www.heroku.com
- Most instructions on Heroku shows how to deploy with Git, but *I do not like it...* we will use a CLI
 - however, you can use whatever you like...
- Install *Heroku CLI*, which allows you to interact with Heroku from command line
- On the web interface, create an “app” with a name of your choice. In these slides, I will use “*pg6300-c4*”
 - as names are unique, you will need to choose a different name

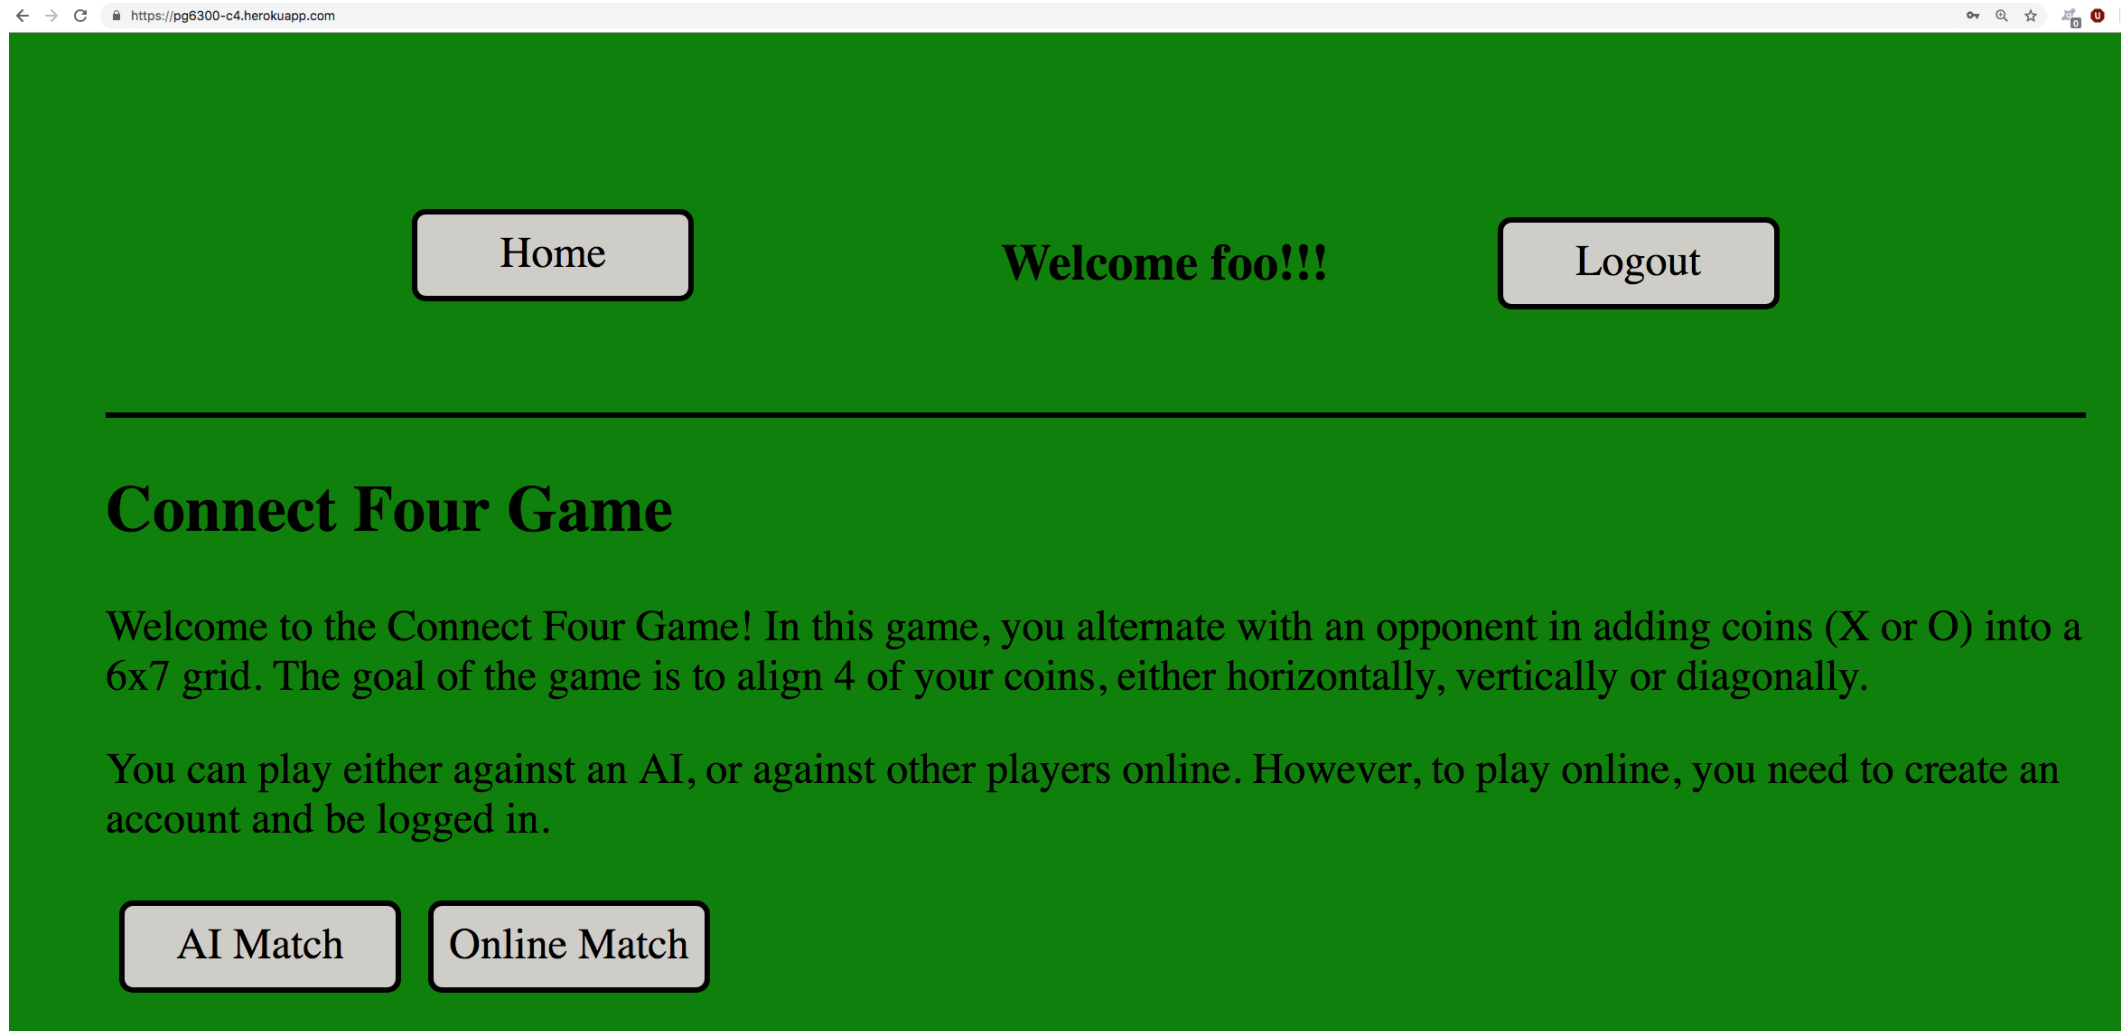
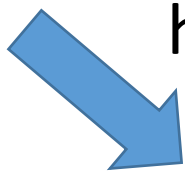
From Command Line (CLI)

- **heroku plugins:install heroku-builds**
 - need to be run only once, to install the “*builds*” plugin
- **heroku login**
 - will setup credential for the other commands.
 - note: if using Windows, this might not work on GitBash, and need to do this command once from a regular Terminal
- **heroku builds:create -a pg6300-c4**
 - zip all your files in current folder, and deploy them in the app
 - note: use “*.gitignore*” to specify what to exclude

Settings

- Can have extra settings in “*engines*” under *package.json*, but not compulsory
 - eg., version of Node and YARN to use to run/build your app
- Your app **MUST** bind to a port specified by **process.env.PORT**
 - otherwise, Heroku will not know how to reverse-proxy to it
- Your app will be automatically built by Heroku with “*yarn build*”, and started with “*yarn start*”

<https://pg6300-c4.herokuapp.com/>



Cloud and WebSockets

- WebSockets are expensive resources
- Free-tier cloud options might allow only small number of WS
- Might shut them down automatically if inactive for even short period, e.g. 30-60 seconds
 - could implement some auto-reconnect when sockets are forcibly closed...

Databases

Why Databases?

- Need to store your data somewhere
- If in memory, lose everything as soon as the app restarts
- For a full app, you need a database
 - cloud providers like Heroku give you databases as well
- We are not going to see databases in this course, but need to briefly discuss them

The 3 Rules of Choosing a Database

1. New project or unsure what to do? **Choose Postgres**
2. If you are already using *MySQL* and migration to *Postgres* would be too expensive, can stick with *MySQL*
3. If you have a long experience with databases, know exactly what you are doing, and can measure objectively the performance benefits of different tradeoffs compared to just using *Postgres*, then, *and only then*, choose best database for the *specific* problem you are facing

Example: *MySQL*

- Open source, but own (and mainly developed) by *Oracle*... and let's not forget that one of its main commercial products is *Oracle Database*...
 - so, yes, in theory those 2 databases are competitors...
- For most use cases, *MySQL* is on par with *Postgres*, but usually slower at adding new advanced features
 - eg support for NoSQL features like JSON data type, or SQL compliance

Example: *MongoDB*

- Most famous *NoSQL* database
 - very, very popular in tutorials... especially in NodeJS
- Meant for *documents*, not for data with *relations*
 - Usually documents are in JSON format, where the only relations are hierarchical, eg nested objects
- Can be *fast* and *easy* to set up...
- ... but you need to *sacrifice ACID* for it...
 - eg, when you “save” some data, can be just cached, and not actually saved...
 - ACID transactions added in v4.0, in 2018...
- *Postgres/MySQL* can save JSON fields, and be very fast at it
 - eg, in 2014, Postgres was actually faster than MongoDB in benchmarks at dealing with JSON

MongoDB Cont.

- Might start with JSON *documents*... but then one day you need to add relations between data: *you are screwed*
 - “screwed” meaning ending up implementing JOINS at application level, which is a nightmare and very inefficient... and/or duplicate data, which need to be kept always in sync...
- Or even worse, choosing *MongoDB* even when you deal with relational data, just because of *hype*...
- ...or when you do not really deal with the amount of data of Google/Amazon/etc...

But... MongoDB is “Web Scale”!

- <https://www.youtube.com/watch?v=b2F-DItXtZs>
- <http://www.sarahmei.com/blog/2013/11/11/why-you-should-never-use-mongodb/>
- **echo "MongoDB is Web Scale!" > /dev/null**
- Note: video is from 2010. At that time MongoDB was total “*rubbish*”. Today is better
 - eg ACID transactions added in 2018

